

데이터 다루기

판다스를 활용한 데이터 조사

```
In [3]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# 피마 인디언 당뇨병 데이터셋을 불러옵니다.
df = pd.read_csv('./data/pima-indians-diabetes3.csv')
```

```
In [4]: # 처음 5줄을 봅니다.
df.head(5)
```

```
Out[4]:
```

	pregnant	plasma	pressure	thickness	insulin	bmi	pedigree	age	diabetes
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [5]: # 정상과 당뇨 환자가 각각 몇 명씩인지 조사해 봅니다.
df["diabetes"].value_counts()
```

```
Out[5]: diabetes
0      500
1      268
Name: count, dtype: int64
```

```
In [6]: # 각 정보별 특징을 좀 더 자세히 출력합니다.  
df.describe()
```

```
Out[6]:
```

	pregnant	plasma	pressure	thickness	insulin	bmi	pedigree	age	diabetes
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000


```
In [7]: # 각 항목이 어느 정도의 상관 관계를 가지고 있는지 알아봅니다.  
df.corr()
```


Out[7]:

	pregnant	plasma	pressure	thickness	insulin	bmi	pedigree	age	diabetes
pregnant	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0.544341	0.221898
plasma	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514	0.466581
pressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528	0.065068
thickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970	0.074752
insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163	0.130548
bmi	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0.036242	0.292695
pedigree	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0.033561	0.173844
age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.033561	1.000000	0.238356
diabetes	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.173844	0.238356	1.000000

```
In [8]: # 데이터 간의 상관 관계를 그래프로 표현해 봅니다.
colormap = plt.cm.gist_heat # 그래프의 색상 구성을 정합니다.
plt.figure(figsize=(12,12)) # 그래프의 크기를 정합니다.

# 그래프의 속성을 결정합니다. vmax의 값을 0.5로 지정해 0.5에 가까울수록 밝은색으로 표시되게 합니다.
sns.heatmap(df.corr(),linewidths=0.1,vmax=0.5, cmap=colormap,
            linecolor='white', annot=True)
plt.show()
```






4. 중요한 데이터 추출하기

```
In [10]: import warnings
warnings.filterwarnings("ignore")


# plasma를 기준으로 각각 정상과 당뇨가 어느 정도 비율로 분포하는지 살펴봅니다.
plt.hist(x=[df.plasma[df.diabetes==0], df.plasma[df.diabetes==1]], bins=30,
         histtype='barstacked', label=['normal', 'diabetes'])
plt.legend()
```

Out[10]: <matplotlib.legend.Legend at 0x1a32dfe6ab0>



```
In [11]: # BMI를 기준으로 각각 정상과 당뇨가 어느 정도 비율로 분포하는지 살펴봅니다.  
plt.hist(x=[df.bmi[df.diabetes==0], df.bmi[df.diabetes==1]], bins=30,  
         histtype='barstacked', label=['normal', 'diabetes'])  
plt.legend()
```

```
Out[11]: <matplotlib.legend.Legend at 0x1a32d277a70>
```



5. 피마 인디언 당뇨병 예측 실행

```
In [13]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

# pandas 라이브러리를 불러옵니다.
import pandas as pd

# 피마 인디언 당뇨병 데이터셋을 불러옵니다.
df = pd.read_csv('./data/pima-indians-diabetes3.csv')
```

```
In [14]: # 세부 정보를 x로 지정합니다.
x = df.iloc[:,0:8]
```

```
# 당뇨병 여부를 y로 지정합니다.  
y = df.iloc[:,8]
```

```
In [15]: # 모델을 설정합니다.  
model = Sequential()  
model.add(Dense(12, input_dim=8, activation='relu', name='Dense_1'))  
model.add(Dense(8, activation='relu', name='Dense_2'))  
model.add(Dense(1, activation='sigmoid', name='Dense_3'))  
model.summary()
```

Model: "sequential"
















Layer (type)	Output Shape	Param #
Dense_1 (Dense)	(None, 12)	108
Dense_2 (Dense)	(None, 8)	104
Dense_3 (Dense)	(None, 1)	9


Total params: 221 (884.00 B)


Trainable params: 221 (884.00 B)


Non-trainable params: 0 (0.00 B)


```
In [16]: # 모델을 컴파일합니다.  
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])  
  
# 모델을 실행합니다.  
history=model.fit(X, y, epochs=100, batch_size=5)
```



Epoch 1/100
154/154  1s 927us/step - accuracy: 0.5365 - loss: 2.8902
Epoch 2/100
154/154  0s 910us/step - accuracy: 0.6390 - loss: 1.1825
Epoch 3/100
154/154  0s 787us/step - accuracy: 0.6601 - loss: 0.9206
Epoch 4/100
154/154  0s 791us/step - accuracy: 0.5786 - loss: 1.0341
Epoch 5/100
154/154  0s 821us/step - accuracy: 0.6447 - loss: 0.8250
Epoch 6/100
154/154  0s 758us/step - accuracy: 0.6410 - loss: 0.8132
Epoch 7/100
154/154  0s 803us/step - accuracy: 0.6430 - loss: 0.8093
Epoch 8/100
154/154  0s 776us/step - accuracy: 0.6580 - loss: 0.7157
Epoch 9/100
154/154  0s 777us/step - accuracy: 0.6203 - loss: 0.8318
Epoch 10/100
154/154  0s 783us/step - accuracy: 0.6725 - loss: 0.6819
Epoch 11/100
154/154  0s 861us/step - accuracy: 0.6297 - loss: 0.7660
Epoch 12/100
154/154  0s 821us/step - accuracy: 0.6571 - loss: 0.6649
Epoch 13/100
154/154  0s 776us/step - accuracy: 0.6630 - loss: 0.6808
Epoch 14/100
154/154  0s 797us/step - accuracy: 0.6263 - loss: 0.8030
Epoch 15/100
154/154  0s 786us/step - accuracy: 0.6719 - loss: 0.6374
Epoch 16/100
154/154  0s 924us/step - accuracy: 0.6564 - loss: 0.7278
Epoch 17/100
154/154  0s 781us/step - accuracy: 0.6711 - loss: 0.6465
Epoch 18/100
154/154  0s 812us/step - accuracy: 0.7208 - loss: 0.6101
Epoch 19/100
154/154  0s 795us/step - accuracy: 0.7027 - loss: 0.6048
Epoch 20/100
154/154  0s 754us/step - accuracy: 0.7143 - loss: 0.6019
Epoch 21/100


154/154  0s 1ms/step - accuracy: 0.6563 - loss: 0.6568
Epoch 22/100


154/154  0s 1ms/step - accuracy: 0.6714 - loss: 0.6535
Epoch 23/100


154/154  0s 781us/step - accuracy: 0.6967 - loss: 0.6017
Epoch 24/100


154/154  0s 786us/step - accuracy: 0.6971 - loss: 0.6470
Epoch 25/100


154/154  0s 808us/step - accuracy: 0.7067 - loss: 0.5890
Epoch 26/100


154/154  0s 844us/step - accuracy: 0.6882 - loss: 0.5977
Epoch 27/100


154/154  0s 827us/step - accuracy: 0.6844 - loss: 0.6217
Epoch 28/100


154/154  0s 787us/step - accuracy: 0.6460 - loss: 0.6364
Epoch 29/100


154/154  0s 794us/step - accuracy: 0.6626 - loss: 0.5985
Epoch 30/100


154/154  0s 820us/step - accuracy: 0.6425 - loss: 0.6307
Epoch 31/100


154/154  0s 852us/step - accuracy: 0.7025 - loss: 0.5741
Epoch 32/100


154/154  0s 806us/step - accuracy: 0.7101 - loss: 0.5915
Epoch 33/100


154/154  0s 923us/step - accuracy: 0.6734 - loss: 0.6301
Epoch 34/100


154/154  0s 799us/step - accuracy: 0.6752 - loss: 0.6221
Epoch 35/100


154/154  0s 782us/step - accuracy: 0.7178 - loss: 0.6084
Epoch 36/100


154/154  0s 835us/step - accuracy: 0.7075 - loss: 0.5971
Epoch 37/100

154/154  0s 804us/step - accuracy: 0.7123 - loss: 0.5593
Epoch 38/100


154/154  0s 774us/step - accuracy: 0.7354 - loss: 0.5713
Epoch 39/100


154/154  0s 795us/step - accuracy: 0.7292 - loss: 0.5929
Epoch 40/100


154/154  0s 780us/step - accuracy: 0.6956 - loss: 0.5879
Epoch 41/100


154/154  0s 824us/step - accuracy: 0.6871 - loss: 0.5982


Epoch 42/100
154/154  0s 857us/step - accuracy: 0.6635 - loss: 0.6448
Epoch 43/100
154/154  0s 808us/step - accuracy: 0.6953 - loss: 0.5737
Epoch 44/100
154/154  0s 764us/step - accuracy: 0.7110 - loss: 0.5824
Epoch 45/100
154/154  0s 1ms/step - accuracy: 0.7511 - loss: 0.5404
Epoch 46/100
154/154  0s 808us/step - accuracy: 0.7206 - loss: 0.5412
Epoch 47/100
154/154  0s 911us/step - accuracy: 0.7112 - loss: 0.6021
Epoch 48/100
154/154  0s 774us/step - accuracy: 0.7165 - loss: 0.6058
Epoch 49/100
154/154  0s 1ms/step - accuracy: 0.7298 - loss: 0.5570
Epoch 50/100
154/154  0s 917us/step - accuracy: 0.7156 - loss: 0.5744
Epoch 51/100
154/154  0s 793us/step - accuracy: 0.7437 - loss: 0.5336
Epoch 52/100
154/154  0s 770us/step - accuracy: 0.7479 - loss: 0.5348
Epoch 53/100
154/154  0s 787us/step - accuracy: 0.7158 - loss: 0.5588
Epoch 54/100
154/154  0s 808us/step - accuracy: 0.7224 - loss: 0.5798
Epoch 55/100
154/154  0s 782us/step - accuracy: 0.6888 - loss: 0.6272
Epoch 56/100
154/154  0s 841us/step - accuracy: 0.7124 - loss: 0.5739
Epoch 57/100
154/154  0s 778us/step - accuracy: 0.7117 - loss: 0.5585
Epoch 58/100
154/154  0s 999us/step - accuracy: 0.7218 - loss: 0.5447
Epoch 59/100
154/154  0s 820us/step - accuracy: 0.7059 - loss: 0.5492
Epoch 60/100
154/154  0s 796us/step - accuracy: 0.7444 - loss: 0.5303
Epoch 61/100
154/154  0s 817us/step - accuracy: 0.7544 - loss: 0.5302
Epoch 62/100


154/154  0s 786us/step - accuracy: 0.7428 - loss: 0.5333
Epoch 63/100


154/154  0s 801us/step - accuracy: 0.7481 - loss: 0.5458
Epoch 64/100


154/154  0s 812us/step - accuracy: 0.7243 - loss: 0.5591
Epoch 65/100


154/154  0s 794us/step - accuracy: 0.7775 - loss: 0.5067
Epoch 66/100


154/154  0s 794us/step - accuracy: 0.7456 - loss: 0.5448
Epoch 67/100


154/154  0s 779us/step - accuracy: 0.7501 - loss: 0.5695
Epoch 68/100


154/154  0s 786us/step - accuracy: 0.7225 - loss: 0.5709
Epoch 69/100


154/154  0s 780us/step - accuracy: 0.7596 - loss: 0.5278
Epoch 70/100


154/154  0s 792us/step - accuracy: 0.7349 - loss: 0.5372
Epoch 71/100

154/154  0s 767us/step - accuracy: 0.7429 - loss: 0.5434
Epoch 72/100


154/154  0s 794us/step - accuracy: 0.7543 - loss: 0.5252
Epoch 73/100

154/154  0s 765us/step - accuracy: 0.7177 - loss: 0.5613
Epoch 74/100


154/154  0s 981us/step - accuracy: 0.7385 - loss: 0.5531
Epoch 75/100


154/154  0s 854us/step - accuracy: 0.7436 - loss: 0.5294
Epoch 76/100


154/154  0s 855us/step - accuracy: 0.7261 - loss: 0.5206
Epoch 77/100


154/154  0s 778us/step - accuracy: 0.7275 - loss: 0.5898
Epoch 78/100

154/154  0s 807us/step - accuracy: 0.7078 - loss: 0.5381
Epoch 79/100

154/154  0s 821us/step - accuracy: 0.7220 - loss: 0.5519
Epoch 80/100

154/154  0s 1ms/step - accuracy: 0.7430 - loss: 0.5291
Epoch 81/100

154/154  0s 842us/step - accuracy: 0.7625 - loss: 0.5137
Epoch 82/100

154/154  0s 753us/step - accuracy: 0.7314 - loss: 0.5326

Epoch 83/100
154/154 ————— 0s 779us/step - accuracy: 0.7763 - loss: 0.5197
Epoch 84/100
154/154 ————— 0s 775us/step - accuracy: 0.7469 - loss: 0.5179
Epoch 85/100
154/154 ————— 0s 767us/step - accuracy: 0.7444 - loss: 0.5158
Epoch 86/100
154/154 ————— 0s 947us/step - accuracy: 0.7056 - loss: 0.5737
Epoch 87/100
154/154 ————— 0s 763us/step - accuracy: 0.7559 - loss: 0.5398
Epoch 88/100
154/154 ————— 0s 827us/step - accuracy: 0.7032 - loss: 0.5613
Epoch 89/100
154/154 ————— 0s 764us/step - accuracy: 0.7651 - loss: 0.4987
Epoch 90/100
154/154 ————— 0s 797us/step - accuracy: 0.7281 - loss: 0.5502
Epoch 91/100
154/154 ————— 0s 757us/step - accuracy: 0.7357 - loss: 0.5226
Epoch 92/100
154/154 ————— 0s 747us/step - accuracy: 0.7396 - loss: 0.5375
Epoch 93/100
154/154 ————— 0s 775us/step - accuracy: 0.7392 - loss: 0.5304
Epoch 94/100
154/154 ————— 0s 778us/step - accuracy: 0.7403 - loss: 0.5225
Epoch 95/100
154/154 ————— 0s 767us/step - accuracy: 0.7700 - loss: 0.4906
Epoch 96/100
154/154 ————— 0s 760us/step - accuracy: 0.7543 - loss: 0.5300
Epoch 97/100
154/154 ————— 0s 768us/step - accuracy: 0.7413 - loss: 0.5390
Epoch 98/100
154/154 ————— 0s 767us/step - accuracy: 0.6862 - loss: 0.5992
Epoch 99/100
154/154 ————— 0s 787us/step - accuracy: 0.7442 - loss: 0.5505
Epoch 100/100
154/154 ————— 0s 767us/step - accuracy: 0.7602 - loss: 0.4981

```
In [17]: print("\n loss: %.4f" % (model.evaluate(X, y)[0]))
```

24/24  0s 1ms/step - accuracy: 0.7416 - loss: 0.5021

loss: 0.4776

```
In [18]: print("\n Accuracy: %.4f" % (model.evaluate(X, y)[1]))
```

24/24  0s 1ms/step - accuracy: 0.7416 - loss: 0.5021

Accuracy: 0.7708