

세상에 없는 얼굴 GAN, 오토인코더

3. 적대적 신경망 실행하기

실습: GAN 모델 만들기

```
In [4]: from tensorflow.keras.datasets import mnist
from tensorflow.keras.layers import Input, Dense, Reshape, Flatten, Dropout
from tensorflow.keras.layers import BatchNormalization, Activation, LeakyReLU, UpSampling2D, Conv2D
from tensorflow.keras.models import Sequential, Model

import numpy as np
import matplotlib.pyplot as plt

# 생성자 모델을 만듭니다.
generator = Sequential()
generator.add(Dense(128*7*7, input_dim=100, activation=LeakyReLU(0.2)))
generator.add(BatchNormalization())
generator.add(Reshape((7, 7, 128)))
generator.add(UpSampling2D())
generator.add(Conv2D(64, kernel_size=5, padding='same'))
generator.add(BatchNormalization())
generator.add(Activation(LeakyReLU(0.2)))
generator.add(UpSampling2D())
generator.add(Conv2D(1, kernel_size=5, padding='same', activation='tanh'))

# 판별자 모델을 만듭니다.
discriminator = Sequential()
discriminator.add(Conv2D(64, kernel_size=5, strides=2, input_shape=(28,28,1), padding="same"))
discriminator.add(Activation(LeakyReLU(0.2)))
discriminator.add(Dropout(0.3))
discriminator.add(Conv2D(128, kernel_size=5, strides=2, padding="same"))
discriminator.add(Activation(LeakyReLU(0.2)))
discriminator.add(Dropout(0.3))
```

```

discriminator.add(Flatten())
discriminator.add(Dense(1, activation='sigmoid'))
discriminator.compile(loss='binary_crossentropy', optimizer='adam')
discriminator.trainable = False

# 생성자와 판별자 모델을 연결시키는 gan 모델을 만듭니다.
ginput = Input(shape=(100,))
dis_output = discriminator(generator(ginput))
gan = Model(ginput, dis_output)
gan.compile(loss='binary_crossentropy', optimizer='adam')
gan.summary()

# 신경망을 실행시키는 함수를 만듭니다.
def gan_train(epoch, batch_size, saving_interval):

    # MNIST 데이터를 불러옵니다.

    (X_train, _), (_, _) = mnist.load_data() # 앞서 불러온 적 있는 MNIST를 다시 이용합니다.
    #단, 테스트 과정은 필요 없고 이미지만 사용할 것이기 때문에 X_train만 불러왔습니다.
    X_train = X_train.reshape(X_train.shape[0], 28, 28, 1).astype('float32')
    X_train = (X_train - 127.5) / 127.5 # 픽셀 값은 0에서 255 사이의 값입니다.
    #이전에 255로 나누어 줄때는 이를 0~1 사이의 값으로 바꾸었던 것인데,
    #여기서는 127.5를 빼준 뒤 127.5로 나누어 줌으로 인해 -1에서 1사이의 값으로 바뀌게 됩니다.
    # X_train.shape, Y_train.shape, X_test.shape, Y_test.shape

    true = np.ones((batch_size, 1))
    fake = np.zeros((batch_size, 1))

    for i in range(epoch):
        # 실제 데이터를 판별자에 입력하는 부분입니다.
        idx = np.random.randint(0, X_train.shape[0], batch_size)
        imgs = X_train[idx]
        d_loss_real = discriminator.train_on_batch(imgs, true)

        # 가상 이미지를 판별자에 입력하는 부분입니다.
        noise = np.random.normal(0, 1, (batch_size, 100))
        gen_imgs = generator.predict(noise, verbose=0)
        d_loss_fake = discriminator.train_on_batch(gen_imgs, fake)

        # 판별자와 생성자의 오차를 계산합니다.
        d_loss = 0.5 * np.add(d_loss_real, d_loss_fake)

```

```

g_loss = gan.train_on_batch(noise, true)
if i % 100 == 0:
    print('epoch:%d' % i, ' d_loss:%.4f' % d_loss, ' g_loss:%.4f' % g_loss)

# 이 부분은 중간 과정을 이미지로 저장해 주는 부분입니다. 이 장의 주요 내용과 관련이 없어
# 소스 코드만 첨부합니다. 만들어진 이미지들은 gan_images 폴더에 저장됩니다.
if i % saving_interval == 0:
    #r, c = 5, 5
    noise = np.random.normal(0, 1, (25, 100))
    gen_imgs = generator.predict(noise)

    # Rescale images 0 - 1
    gen_imgs = 0.5 * gen_imgs + 0.5

    fig, axs = plt.subplots(5, 5)
    count = 0
    for j in range(5):
        for k in range(5):
            axs[j, k].imshow(gen_imgs[count, :, :, 0], cmap='gray')
            axs[j, k].axis('off')
            count += 1
    fig.savefig("./gan_mnist_%d.png" % i)

gan_train(2001, 32, 200) # 2000번 반복되고, 배치 사이즈는 32, 200번마다 결과가 저장되게 하였습니다.

```

C:\Users\user\AppData\Roaming\Python\Python312\site-packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

C:\Users\user\AppData\Roaming\Python\Python312\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "functional_17"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 100)	0
sequential (Sequential)	(None, 28, 28, 1)	865,281
sequential_1 (Sequential)	(None, 1)	212,865


Total params: 1,078,146 (4.11 MB)


Trainable params: 852,609 (3.25 MB)


Non-trainable params: 225,537 (881.00 KB)

C:\Users\user\AppData\Roaming\Python\Python312\site-packages\keras\src\backend\tensorflow\trainer.py:82: UserWarning: The model does not have any trainable weights.
warnings.warn("The model does not have any trainable weights.")


```
epoch:0 d_loss:0.6919 g_loss:0.6950
1/1 ————— 0s 122ms/step
epoch:100 d_loss:0.7522 g_loss:0.5738
epoch:200 d_loss:0.8225 g_loss:0.4942
1/1 ————— 0s 28ms/step
epoch:300 d_loss:0.8872 g_loss:0.4327
epoch:400 d_loss:0.9380 g_loss:0.3897
1/1 ————— 0s 36ms/step
epoch:500 d_loss:0.9779 g_loss:0.3581
epoch:600 d_loss:1.0103 g_loss:0.3339
1/1 ————— 0s 59ms/step
epoch:700 d_loss:1.0364 g_loss:0.3151
epoch:800 d_loss:1.0584 g_loss:0.2999
1/1 ————— 0s 30ms/step
epoch:900 d_loss:1.0764 g_loss:0.2876
epoch:1000 d_loss:1.0914 g_loss:0.2776
1/1 ————— 0s 40ms/step
epoch:1100 d_loss:1.1041 g_loss:0.2692
epoch:1200 d_loss:1.1148 g_loss:0.2621
1/1 ————— 0s 38ms/step
epoch:1300 d_loss:1.1242 g_loss:0.2561
epoch:1400 d_loss:1.1323 g_loss:0.2508
1/1 ————— 0s 30ms/step
epoch:1500 d_loss:1.1393 g_loss:0.2463
epoch:1600 d_loss:1.1454 g_loss:0.2424
1/1 ————— 0s 29ms/step
epoch:1700 d_loss:1.1510 g_loss:0.2388
epoch:1800 d_loss:1.1560 g_loss:0.2357
1/1 ————— 0s 28ms/step
epoch:1900 d_loss:1.1603 g_loss:0.2329
epoch:2000 d_loss:1.1644 g_loss:0.2303
1/1 ————— 0s 29ms/step
```





























4. 이미지의 특징을 추출하는 오토인코더

실습: 오토인코더 실습하기

```
In [7]: from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import Input, Dense, Conv2D, MaxPooling2D, UpSampling2D, Flatten, Reshape

import matplotlib.pyplot as plt
import numpy as np

# MNIST 데이터셋을 불러옵니다.

(X_train, _), (X_test, _) = mnist.load_data()
```



```
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1).astype('float32') / 255
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1).astype('float32') / 255

# 생성자 모델을 만듭니다.
autoencoder = Sequential()

# 인코딩 부분입니다.
autoencoder.add(Conv2D(16, kernel_size=3, padding='same', input_shape=(28,28,1), activation='relu'))
autoencoder.add(MaxPooling2D(pool_size=2, padding='same'))
autoencoder.add(Conv2D(8, kernel_size=3, activation='relu', padding='same'))
autoencoder.add(MaxPooling2D(pool_size=2, padding='same'))
autoencoder.add(Conv2D(8, kernel_size=3, strides=2, padding='same', activation='relu'))

# 디코딩 부분입니다.
autoencoder.add(Conv2D(8, kernel_size=3, padding='same', activation='relu'))
autoencoder.add(UpSampling2D())
autoencoder.add(Conv2D(8, kernel_size=3, padding='same', activation='relu'))
autoencoder.add(UpSampling2D())
autoencoder.add(Conv2D(16, kernel_size=3, activation='relu'))
autoencoder.add(UpSampling2D())
autoencoder.add(Conv2D(1, kernel_size=3, padding='same', activation='sigmoid'))

# 전체 구조를 확인합니다.
autoencoder.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 28, 28, 16)	160
max_pooling2d (MaxPooling2D)	(None, 14, 14, 16)	0
conv2d_5 (Conv2D)	(None, 14, 14, 8)	1,160
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 8)	0
conv2d_6 (Conv2D)	(None, 4, 4, 8)	584
conv2d_7 (Conv2D)	(None, 4, 4, 8)	584
up_sampling2d_2 (UpSampling2D)	(None, 8, 8, 8)	0
conv2d_8 (Conv2D)	(None, 8, 8, 8)	584
up_sampling2d_3 (UpSampling2D)	(None, 16, 16, 8)	0
conv2d_9 (Conv2D)	(None, 14, 14, 16)	1,168
up_sampling2d_4 (UpSampling2D)	(None, 28, 28, 16)	0
conv2d_10 (Conv2D)	(None, 28, 28, 1)	145

Total params: 4,385 (17.13 KB)

Trainable params: 4,385 (17.13 KB)

Non-trainable params: 0 (0.00 B)

```
In [8]: # 컴파일 및 학습을 하는 부분입니다.
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
autoencoder.fit(X_train, X_train, epochs=50, batch_size=128, validation_data=(X_test, X_test))

# 학습된 결과를 출력하는 부분입니다.
random_test = np.random.randint(X_test.shape[0], size=5) # 테스트할 이미지를 랜덤하게 불러옵니다.
ae_imgs = autoencoder.predict(X_test)                    # 앞서 만든 오토인코더 모델에 집어 넣습니다.
```

```
plt.figure(figsize=(7, 2)) # 출력될 이미지의 크기를 정합니다.

for i, image_idx in enumerate(random_test): # 랜덤하게 뽑은 이미지를 차례로 나열합니다.
    ax = plt.subplot(2, 7, i + 1)
    plt.imshow(X_test[image_idx].reshape(28, 28)) # 테스트할 이미지를 먼저 그대로 보여줍니다.
    ax.axis('off')
    ax = plt.subplot(2, 7, 7 + i + 1)
    plt.imshow(ae_imgs[image_idx].reshape(28, 28)) # 오토인코딩 결과를 다음열에 출력합니다.
    ax.axis('off')
plt.show()
```

Epoch 1/50

469/469  5s 8ms/step - loss: 0.3217 - val_loss: 0.1399

Epoch 2/50

469/469  4s 8ms/step - loss: 0.1342 - val_loss: 0.1197

Epoch 3/50

469/469  4s 8ms/step - loss: 0.1176 - val_loss: 0.1100

Epoch 4/50

469/469  4s 8ms/step - loss: 0.1095 - val_loss: 0.1048

Epoch 5/50

469/469  4s 8ms/step - loss: 0.1049 - val_loss: 0.1013

Epoch 6/50

469/469  4s 8ms/step - loss: 0.1017 - val_loss: 0.0987

Epoch 7/50

469/469  4s 8ms/step - loss: 0.0995 - val_loss: 0.0970

Epoch 8/50

469/469  4s 8ms/step - loss: 0.0980 - val_loss: 0.0962

Epoch 9/50

469/469  4s 8ms/step - loss: 0.0966 - val_loss: 0.0945

Epoch 10/50

469/469  4s 7ms/step - loss: 0.0955 - val_loss: 0.0935

Epoch 11/50

469/469  4s 8ms/step - loss: 0.0946 - val_loss: 0.0926

Epoch 12/50

469/469  4s 8ms/step - loss: 0.0937 - val_loss: 0.0919

Epoch 13/50

469/469  4s 8ms/step - loss: 0.0930 - val_loss: 0.0911

Epoch 14/50

469/469  4s 8ms/step - loss: 0.0921 - val_loss: 0.0905

Epoch 15/50

469/469  4s 8ms/step - loss: 0.0915 - val_loss: 0.0900

Epoch 16/50

469/469  4s 8ms/step - loss: 0.0911 - val_loss: 0.0893

Epoch 17/50

469/469  4s 8ms/step - loss: 0.0906 - val_loss: 0.0887

Epoch 18/50

469/469  4s 8ms/step - loss: 0.0900 - val_loss: 0.0885


Epoch 19/50


469/469  4s 8ms/step - loss: 0.0892 - val_loss: 0.0878


Epoch 20/50


469/469  4s 8ms/step - loss: 0.0886 - val_loss: 0.0872


Epoch 21/50


469/469  4s 7ms/step - loss: 0.0882 - val_loss: 0.0869
Epoch 22/50


469/469  4s 8ms/step - loss: 0.0880 - val_loss: 0.0865
Epoch 23/50


469/469  4s 8ms/step - loss: 0.0875 - val_loss: 0.0864
Epoch 24/50


469/469  4s 8ms/step - loss: 0.0872 - val_loss: 0.0858
Epoch 25/50


469/469  4s 8ms/step - loss: 0.0867 - val_loss: 0.0857
Epoch 26/50


469/469  4s 8ms/step - loss: 0.0862 - val_loss: 0.0850
Epoch 27/50


469/469  4s 8ms/step - loss: 0.0860 - val_loss: 0.0848
Epoch 28/50


469/469  4s 8ms/step - loss: 0.0855 - val_loss: 0.0846
Epoch 29/50


469/469  4s 8ms/step - loss: 0.0855 - val_loss: 0.0843
Epoch 30/50


469/469  4s 8ms/step - loss: 0.0854 - val_loss: 0.0840
Epoch 31/50


469/469  4s 8ms/step - loss: 0.0849 - val_loss: 0.0842
Epoch 32/50


469/469  4s 8ms/step - loss: 0.0849 - val_loss: 0.0837
Epoch 33/50


469/469  4s 8ms/step - loss: 0.0846 - val_loss: 0.0834
Epoch 34/50


469/469  4s 8ms/step - loss: 0.0846 - val_loss: 0.0832
Epoch 35/50


469/469  4s 8ms/step - loss: 0.0842 - val_loss: 0.0830
Epoch 36/50


469/469  4s 8ms/step - loss: 0.0841 - val_loss: 0.0829
Epoch 37/50

469/469  4s 8ms/step - loss: 0.0840 - val_loss: 0.0828
Epoch 38/50

469/469  4s 8ms/step - loss: 0.0839 - val_loss: 0.0825
Epoch 39/50

469/469  4s 8ms/step - loss: 0.0837 - val_loss: 0.0825
Epoch 40/50

469/469  4s 8ms/step - loss: 0.0836 - val_loss: 0.0823
Epoch 41/50

469/469  4s 8ms/step - loss: 0.0835 - val_loss: 0.0822

Epoch 42/50
469/469 ————— 4s 8ms/step - loss: 0.0834 - val_loss: 0.0820
Epoch 43/50
469/469 ————— 4s 8ms/step - loss: 0.0829 - val_loss: 0.0819
Epoch 44/50
469/469 ————— 4s 8ms/step - loss: 0.0831 - val_loss: 0.0817
Epoch 45/50
469/469 ————— 4s 8ms/step - loss: 0.0827 - val_loss: 0.0818
Epoch 46/50
469/469 ————— 4s 8ms/step - loss: 0.0828 - val_loss: 0.0822
Epoch 47/50
469/469 ————— 4s 7ms/step - loss: 0.0825 - val_loss: 0.0813
Epoch 48/50
469/469 ————— 4s 8ms/step - loss: 0.0826 - val_loss: 0.0814
Epoch 49/50
469/469 ————— 4s 8ms/step - loss: 0.0823 - val_loss: 0.0813
Epoch 50/50
469/469 ————— 4s 8ms/step - loss: 0.0824 - val_loss: 0.0812
313/313 ————— 1s 3ms/step



In []: